

Comp 135 Introduction to Machine Learning and Data Mining

Fall 2015

Professor: Roni Khardon

Computer Science
Tufts University

Recall Linear Threshold Units

- The basic model:

$$\text{Output} = f(\sum w_j x_j)$$

- Where f can be one of

$$f = \text{sign}() \text{ Value in } \{-1, 1\}$$

$$f = \text{step}() \text{ Value in } \{0, 1\}$$

$$\sigma(a) = \frac{1}{1+e^{-a}}$$

$$p(f = 1) = \sigma(\sum w_j x_j)$$

comp135

Roni Khardon, Tufts University

- Note: we are focusing on one example and omitting the index i saying that this is the i 'th example to avoid clutter in notation

comp135

Roni Khardon, Tufts University

Linear Sigmoid Units

- Today we will work with units whose output is a real value in $[0,1]$

$$\text{Output} = \hat{y} = \sigma(\sum w_j x_j)$$

$$\sigma(a) = \frac{1}{1+e^{-a}}$$

- This conveniently satisfies

$$\sigma'(a) = \frac{-(-e^{-a})}{(1+e^{-a})^2} = \sigma(a)(1 - \sigma(a))$$

comp135

Roni Khardon, Tufts University

Linear Sigmoid Units

- Consider an example (x,y)
- And error function

$$\text{Err} = \frac{1}{2}[y - \hat{y}]^2 = \frac{1}{2}[y - \sigma(\sum_j w_j x_j)]^2$$

- Applying gradient descent

$$w_k = w_k - \eta \frac{\partial \text{Err}}{\partial w_k}$$

- We get the update rule

$$w_k = w_k + \eta(y - \hat{y})\hat{y}(1 - \hat{y})x_k$$

comp135

Roni Khardon, Tufts University


Multi Layer Networks

- Must first develop convenient notation
- This is different from single unit notation
- But it simplifies the exposition of the algorithm that follows

comp135

Roni Khardon, Tufts University

Multi Layer Networks



comp135

Roni Khardon, Tufts University


Multi Layer Networks

- **Must first develop convenient notation**
- Denote input as before by x_1, \dots, x_n
- An internal node is identified by its index i , and its output is x_i
- All internal nodes are x_{n+1}, \dots, x_N
- And the final output is x_N
- The link from unit j to i has weight $w_{j,i}$
- The sum at unit i is $s_i = \sum_j w_{j,i} x_j$
- The output at i is $x_i = \sigma(s_i) = \sigma(\sum w_{j,i} x_j)$

comp135

Roni Khardon, Tufts University

Multi Layer Networks



What does the network predict on this example?

$$\eta = 0.1$$

$$w_{56} = w_{67} = 1$$

$$w_{35} = w_{36} = w_{45} = w_{46} = 0.6$$

$$w_{13} = w_{14} = w_{23} = w_{24} = 1$$

Input example : $(x_1, x_2) = (2, 3)$
 Desired output : $L = 0$

comp135

Roni Khardon, Tufts University

Example

First Step: compute s_i, x_i , and $\sigma'_i = x_i(1 - x_i)$

$$s_3 = 1 * 2 + 1 * 3 = 5 \quad x_3 = \frac{1}{1+e^{-5}} = 0.993 \quad \sigma'_3 = 0.007$$

$$s_4 = 5 \quad x_4 = 0.993 \quad \sigma'_4 = 0.007$$

$$s_5 = 0.6 * 0.993 + 0.6 * 0.993 = 1.192 \quad x_5 = 0.767 \quad \sigma'_5 = 0.179$$

$$s_6 = 1.192 \quad x_6 = 0.767 \quad \sigma'_6 = 0.179$$

$$s_7 = 1 * 0.767 + 1 * 0.767 = 1.534 \quad x_7 = \frac{1}{1+e^{-1.534}} = 0.823 \quad \sigma'_7 = 0.146$$

comp135

Roni Khardon, Tufts University

Multi Layer Networks

- As before we get an example (x, y) .
- x specifies the input units x_1, \dots, x_n
- y is the intended output of x_N
- Nothing is known about intention for middle layers (a.k.a. hidden units)
- Apply same error function
- And gradient descent

comp135

Roni Khardon, Tufts University

Multi Layer Networks

- The error function

$$\text{Err} = \frac{1}{2} [y - x_N]^2$$

- Gradient update:

$$w_{j,i} = w_{j,i} - \eta \frac{\partial \text{Err}}{\partial w_{j,i}}$$

- How can we calculate the gradient for an arbitrary $w_{j,i}$ (at middle or top layer)?

comp135

Roni Khardon, Tufts University

Multi Layer Networks

- The error function

$$\text{Err} = \frac{1}{2}[y - x_N]^2$$

- Gradient update:

$$w_{j,i} = w_{j,i} - \eta \frac{\partial \text{Err}}{\partial w_{j,i}}$$

- Two basic observations:

$$\frac{\partial \text{Err}}{\partial w_{j,i}} = \frac{\partial \text{Err}}{\partial s_i} \frac{\partial s_i}{\partial w_{j,i}}$$

$w_{j,i}$ affects output only through s_i

$$\frac{\partial s_i}{\partial w_{j,i}} = \frac{\partial \sum_j w_{j,i} x_j}{\partial w_{j,i}} = x_j$$

Just a derivative of linear function

comp135

Roni Khardon, Tufts University

Multi Layer Networks

- The error function

$$\text{Err} = \frac{1}{2}[y - x_N]^2$$

- Gradient update:

$$w_{j,i} = w_{j,i} - \eta \frac{\partial \text{Err}}{\partial w_{j,i}}$$

- Two basic observations:

$$\frac{\partial \text{Err}}{\partial w_{j,i}} = \frac{\partial \text{Err}}{\partial s_i} \frac{\partial s_i}{\partial w_{j,i}}$$

$$\frac{\partial \text{Err}}{\partial w_{j,i}} = \frac{\partial \text{Err}}{\partial s_i} \frac{\partial s_i}{\partial w_{j,i}} = \Delta_i x_j$$

$$\frac{\partial s_i}{\partial w_{j,i}} = \frac{\partial \sum_j w_{j,i} x_j}{\partial w_{j,i}} = x_j$$

comp135

Roni Khardon, Tufts University

Backpropagation Algorithm

- A few more steps (on the board) yield the Backpropagation algorithm
- Start by initializing all $w_{j,i}$ to small random values

comp135

Roni Khardon, Tufts University

Backpropagation Algorithm

- Repeat for each example (x,y) :
 - For all i , compute values s_i and x_i by
going forward in network

$$s_i = \sum_j w_{j,i} x_j$$

$$x_i = \sigma(s_i) = \sigma(\sum_j w_{j,i} x_j)$$

- For all i , compute values Δ_i by
going backward in network

$$\Delta_N = -(y - x_N)x_N(1 - x_N)$$

$$\Delta_i = x_i(1 - x_i) \sum_k \Delta_k w_{i,k}$$

- Update all weights

$$w_{j,i} = w_{j,i} - \eta \Delta_i x_j$$

comp135

Roni Khardon, Tufts University

Backpropagation Algorithm

- Algorithm on previous slide updates after each example
- This is known as "stochastic gradient descent" (similar to perceptron)
- The standard Backpropagation algorithm makes multiple iterations over training set: in each iteration it collects the gradients from all examples in the training set and only then makes an update.

comp135

Roni Khardon, Tufts University

Multi Layer Networks




Illustration of Backpropagation

$$\eta = 0.1$$

$$w_{56} = w_{67} = 1$$

$$w_{35} = w_{36} = w_{45} = w_{46} = 0.6$$

$$w_{13} = w_{14} = w_{23} = w_{24} = 1$$

Input example : $(x_1, x_2) = (2, 3)$

Desired output : $L = 0$

comp135

Roni Khardon, Tufts University

Backpropagation Example

First Step: compute s_i, x_i , and $\sigma'_i = x_i(1 - x_i)$

$$s_3 = 1 * 2 + 1 * 3 = 5 \quad x_3 = \frac{1}{1+e^{-5}} = 0.993 \quad \sigma'_3 = 0.007$$

$$s_4 = 5 \quad x_4 = 0.993 \quad \sigma'_4 = 0.007$$

$$s_5 = 0.6 * 0.993 + 0.6 * 0.993 = 1.192 \quad x_5 = 0.767 \quad \sigma'_5 = 0.179$$

$$s_6 = 1.192 \quad x_6 = 0.767 \quad \sigma'_6 = 0.179$$

$$s_7 = 1 * 0.767 + 1 * 0.767 = 1.534 \quad x_7 = \frac{1}{1+e^{-1.534}} = 0.823 \quad \sigma'_6 = 0.146$$

comp135

Roni Khardon, Tufts University

Backpropagation Example

Second Step: compute Δ_i

$$\Delta_7 = -\sigma'_7 * (L - x_7) = -0.146 * (0 - 0.823) = 0.120$$

$$\Delta_5 = \sigma'_5 w_{57} \Delta_7 = 0.179 * 1 * 0.120 = 0.021$$

$$\Delta_6 = \Delta_5$$

$$\Delta_3 = \sigma'_3 [w_{35} \Delta_5 + w_{36} \Delta_6] = 0.000176$$

$$\Delta_4 = \Delta_3$$

comp135

Roni Khardon, Tufts University

Backpropagation Example

Third Step: update weights

$$w_{13} = w_{13} - \eta x_1 \Delta_3 = 1 - 0.1 * 2 * 0.000176 = 0.9999648$$

...

$$w_{35} = w_{35} - \eta x_3 \Delta_5 = 0.6 - 0.1 * 0.993 * 0.021 = 0.5579$$

...


$$w_{57} = w_{57} - \eta x_5 \Delta_7 = 1 - 0.1 * 0.767 * 0.120 = 0.9908$$

...

comp135

Roni Khardon, Tufts University

Multiple Output Nodes



comp135

Roni Khardon, Tufts University

Multiple Output Nodes

- All outputs share the same hidden layer
- Network identifies useful representations that are useful for all outputs
- Exactly same algorithm applies
- Forward pass identical
- Backward pass: each output unit calculates Delta using Δ_N formula

comp135

Roni Khardon, Tufts University

Multi Layer Networks

- Not easy to optimize; the error surface has a lot of local minima
- Solutions:
- Momentum:


$$w_{j,i} = w_{j,i} - \eta \left[\frac{\partial \text{Err}}{\partial w_{j,i}} + \alpha \text{ previous update} \right]$$
- Use multiple restarts and pick one with lowest training set error
- ... many more recent techniques

comp135

Roni Khardon, Tufts University

What does the hidden layer do?

- Example: self-encoders



Learned hidden layer representation:


Input	Hidden Values	Output
10000000	→ .89 .04 .08	→ 10000000
01000000	→ .01 .11 .88	→ 01000000
00100000	→ .01 .97 .27	→ 00100000
00010000	→ .99 .97 .71	→ 00010000
00001000	→ .03 .05 .02	→ 00001000
00000100	→ .22 .99 .99	→ 00000100
00000010	→ .80 .01 .98	→ 00000010
00000001	→ .60 .94 .01	→ 00000001

[Images from Mitchell's textbook]

comp135

Roni Khardon, Tufts University


What does the hidden layer do?



comp135

Roni Khardon, Tufts University

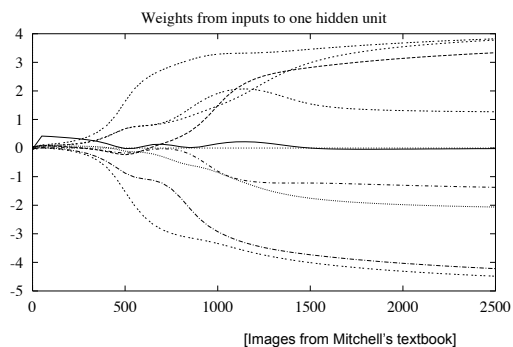
What does the hidden layer do?



comp135

Roni Khardon, Tufts University

What does the hidden layer do?



comp135

Roni Khardon, Tufts University

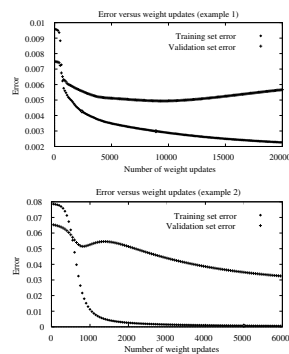
Multi Layer Networks

- How to pick network size (and shape)?
- Similar to model selection in other models
 - cross validation
 - Combine fit + penalty
- How many updates?
 - Overfitting with large number of updates
 - Can do with with large network and moderate number of updates

comp135

Roni Khardon, Tufts University

Multi Layer Networks



Using validation set for stopping criterion per number of updates

comp135

Roni Khardon, Tufts University

Multi Layer Networks

- Renewed interest in **Deep Networks** in last decade
- Several schemes for special network structure and special node functions
- Several schemes for training
- Combination of these ideas with BigData
- Yields
- Impressive improvements in performance in vision and other applications

comp135

Roni Khardon, Tufts University

Convolutional Networks

- Architecture inspired by vision system
- Alternating layers of grid based structures
- Each node calculates local function on patch from previous layer

comp135

Roni Khardon, Tufts University

Convolutional Networks

- Alternate layers of:
- "convolution layer" applies filter to patch from previous layer; weights repeat in all nodes (i.e. same filter)
- "Pooling layer" combines multiple filters of same block
- Followed by fully connected layers

comp135

Roni Khardon, Tufts University

Deep Networks

- Autoencoders: similar to 8-3-8 idea.
- Network fragments can be used to learn one level of internal representations in an unsupervised manner
- Restricted Boltzmann Machines (RBM): a probabilistic model with similar intuitive role
- Stacking these gives a deep network
- Further supervised training of entire model after this step

comp135

Roni Khardon, Tufts University

Deep Networks

- Active area of research
- Still not well understood
- Public interest due to empirical success
- Source of success: Huge data? Network architecture? Training algorithms? Domain specific engineering?

comp135

Roni Khardon, Tufts University

Multi Layer Neural Networks

- Complex representation of functions
- Can be trained with gradient based methods
- But training can be tricky
- Hidden layer "learning representation"
- Recent work on deep networks adds special architecture and/or training procedures

comp135

Roni Khardon, Tufts University